

# An Adaptive Index for Hierarchical Database Systems

Rafael Kallis

BSc Thesis

February 4, 2018



University of  
Zurich UZH

Department of Informatics

## The Workload-Aware Property Index (WAPI):

- Detects frequently updated nodes
- Stops pruning such volatile nodes
- Significantly improves update throughput

## Unproductive Nodes are an unwanted byproduct:

- When the workload changes, volatile nodes cease to be volatile
- They waste space and slow down queries
- They do not contribute to a query match and contain no data


In this thesis we:

- Design and implement two solutions in order to mitigate unproductive nodes
- Analyze the factors impacting the production of unproductive nodes
- Empirically evaluate and compare our two solutions


A Content Management System's (CMS) workload is:

- skewed
- update-heavy
- changing


Skewed workload: small subset of nodes gets frequently updated



No skew




Normal skew



High skew

Changing workload: as time passes, hotspots change

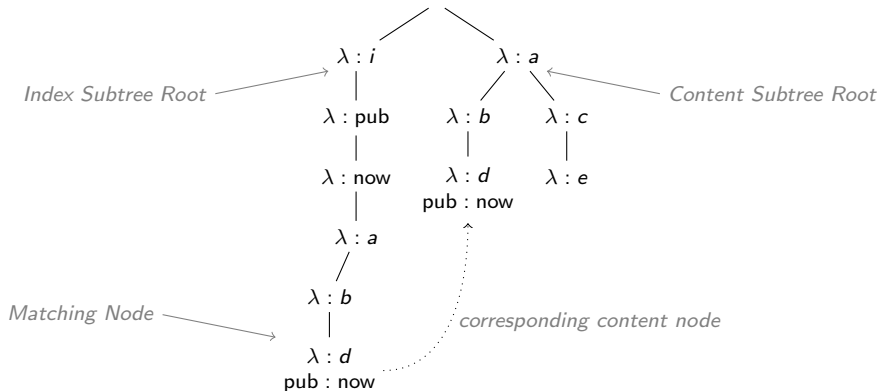


CMSs usually use a job-queuing system that has the noted characteristics



# Workload-Aware Property Index

# Hierarchical Database with WAPI



We mostly executes content-and-structure (CAS) queries [1]. We denote node  $n$ 's property  $k$  as  $n[k]$  and node  $n$ 's descendants as  $desc(n)$ .

### Definition (CAS Query)

Given node  $m$ , property  $k$  and value  $v$ , a CAS query  $Q(k, v, m)$  returns all descendants of  $m$  which have  $k$  set to  $v$ , i.e.,

$$Q(k, v, m) = \{n \mid n \in desc(m) \wedge n[k] = v\}$$

Volatility is the measure which is used by the WAPI in order to distinguish whether to remove a node or not from the index. Wellenzohn et al. [2] propose to look at the recent transactional workload to check whether a node  $n$  is volatile.

## Definition (Volatility Count)

The volatility count  $vol(n)$  of index node  $n$  on database instance  $O_i$ , is the number of times node  $n$  was added or removed from snapshots contained in a Sliding Window of Length  $L$  over history  $H_i$ , i.e.,

$$vol(n) = |\{G^b | G^b \in H_i \wedge t(G^b) \in [t_n - L + 1, t_n] \wedge \exists G^a [ \\ G^a = pre(G^b) \wedge ([n^a \notin N(G^a) \wedge n^b \in N(G^b)] \vee \\ [n^a \in N(G^a) \wedge n^b \notin N(G^b)])]\}|$$

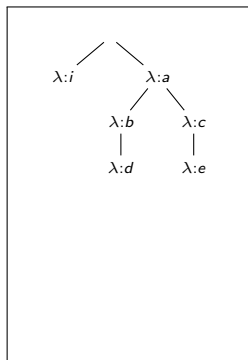
## Definition (Volatile Node)

Index node  $n$  is volatile iff  $n$ 's volatility count is greater or equal than the volatility threshold  $\tau$ , i.e.,

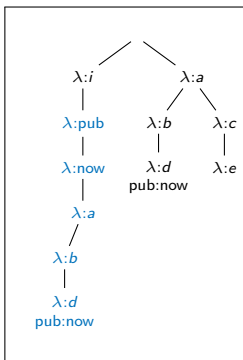
$$\text{volatile}(n) \iff \text{vol}(n) \geq \tau$$

# Index nodes becoming volatile


$$G^0 \xrightarrow{T_1} G^1 \xrightarrow{T_2} G^2$$



Snapshot  $G^0$   
 $t(G^0) = t$



Snapshot  $G^1$   
 $t(G^1) = t + 1$



Snapshot  $G^2$   
 $t(G^2) = t + 2$

$$\tau = 1, L = 2$$

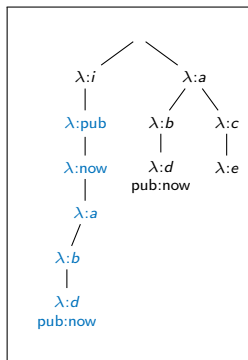
# Unproductive Nodes



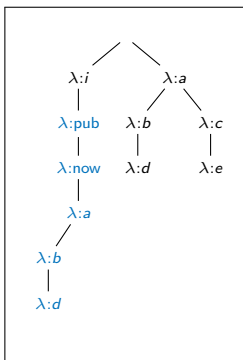
- When nodes are volatile, their volatility count has to be at least  $\tau$
- Time passes and the database workload changes
- Insertions and deletions that increased the volatility count drop out of the sliding window
- If the volatility count drops below threshold  $\tau$ , the node ceases to be volatile
- If the same holds for the node's descendants, we call the node and its descendants unproductive

# Index nodes becoming unproductive


$$G^1 \xrightarrow{T_2} G^2 \xrightarrow{T_3} G^3$$



Snapshot  $G^1$   
 $t(G^1) = t + 1$



Snapshot  $G^2$   
 $t(G^2) = t + 2$



Snapshot  $G^3$   
 $t(G^3) = t + 3$

$\tau = 1, L = 2$

Variables impacting unproductive node production rate:

- Volatility threshold  $\tau$
- Sliding window length  $L$
- Workload skew  $s$
- Update tx per second

Unproductive index node cleaning, we propose:

- Periodic Garbage Collection (GC)
- Query-Time Pruning (QTP)

## Periodic Garbage Collection (GC)

# Periodic GC

Main idea:

- Background process
- Periodically traverse index subtree
- Prune any visited unproductive node

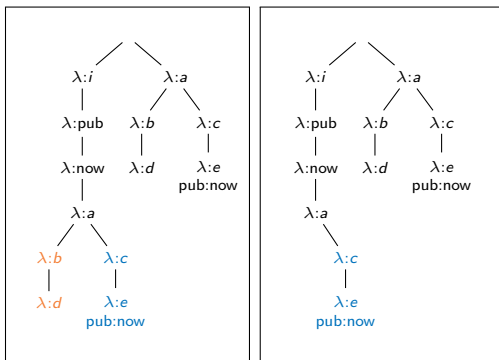
# Periodic GC

## Algorithm: GarbageCollect

```
for node  $n \in \text{desc}(/i)$  in postorder tree walk do  
  if  $\text{chd}(n) = \emptyset \wedge \neg \text{matching}(n) \wedge \neg \text{volatile}(n)$  then  
    delete node  $n$ 
```

# Periodic GC

$$G^3 \xrightarrow{T_4} G^4$$



Snapshot  $G^3$   
 $t(G^3) = t + 3$

Snapshot  $G^4$   
 $t(G^4) = t + 4$

$$\tau = 1, L = 2$$



## Query-Time Pruning (QTP)

# Query-Time Pruning

Main idea:

- Prune unproductive nodes during query execution
- Piggypacking on query execution
- Adds overhead on query runtime

# Query-Time Pruning

**Algorithm:** QueryQTP

**Data:** Query  $Q(k, v, m)$ , where  $k$  is a property,  $v$  a value and  $m (= / \lambda_1 / \dots / \lambda_d)$  a content node's path.

**Result:** A set of nodes satisfying  $Q(k, v, m)$

$r \leftarrow \emptyset$

**for** node  $n \in \text{desc}(/i/k/v/\lambda_1/\dots/\lambda_d)$  *in postorder tree walk*

**do**

**if**  $\text{matching}(n)$  **then**

$r \leftarrow r \cup \{n\}$


**else if**  $\text{chd}(n) = \emptyset \wedge \neg \text{volatile}(n)$  **then**

        delete node  $n$


**return**  $r$

# Query-Time Pruning

$$G^5 \xrightarrow{T_6} G^6$$



Snapshot  $G^5$





Snapshot  $G^6$

$\tau = 1, L = 2, Q(\text{pub}, \text{now}, /a/b)$

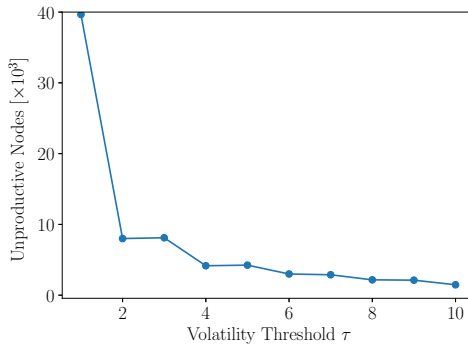
## Experimental Evaluation



# Impact on Query Runtime



# Volatility threshold $\tau$






# Volatility threshold $\tau$

- $\tau$  increases  $\implies$  nodes are less likely to become volatile
- Fewer volatile nodes  $\implies$  fewer unproductive nodes
- Power law relationship between #unproductive nodes and  $\tau$

# Workload skew $s$



# Workload skew $s$

- very high  $s \implies$  very small hotspot
- very small  $s$  (uniform)  $\implies$  no hotspot
- very small/no hotspot  $\implies$  nodes are less likely to become volatile and later on unproductive

# References



MATHIS, C., HÄRDER, T., SCHMIDT, K., AND BÄCHLE, S.  
XML indexing and storage: fulfilling the wish list.  
*Computer Science - R&D* 30, 1 (2015), 51–68.



WELLENZOHN, K., BÖHLEN, M., HELMER, S., REUTEGGER, M., AND SAKR, S.  
A Workload-Aware Index for Tree-Structured Data.  
To be published.