

An Adaptive Index for Hierarchical Database Systems

Rafael Kallis

BSc Thesis

February 14, 2018



University of
Zurich UZH

Department of Informatics

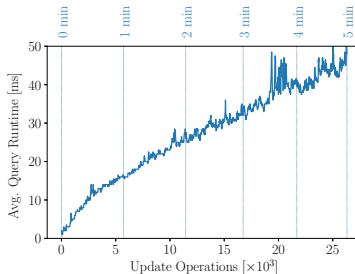
- Skewed and update-heavy workloads trigger repeated structural index updates over a small subset of nodes to the index
- Informally, a frequently added or removed node is called volatile
- Volatile nodes decrease index update performance

The Workload-Aware Property Index (WAPI):

- Detects such volatile nodes
- Stops pruning volatile nodes
- Significantly improves update throughput

Unproductive nodes are an unwanted byproduct:

- When the workload changes, volatile nodes cease to be volatile
- They do not contribute to a query match and contain no data
- They waste space and slow down queries over time




In this thesis we:

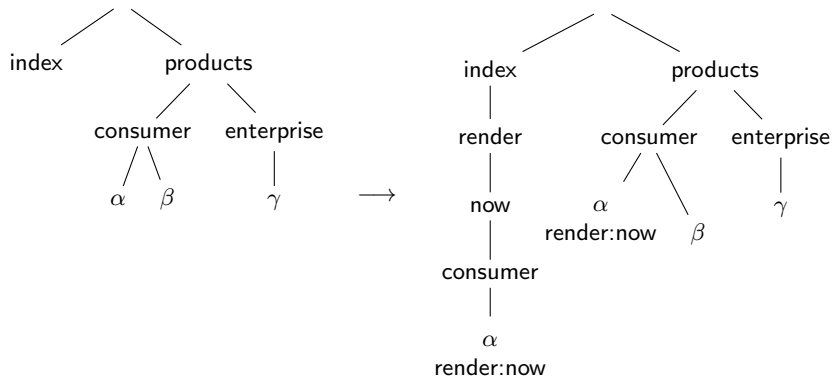
- Design and implement two solutions in order to mitigate unproductive nodes
- Analyze factors impacting the production of unproductive nodes
- Empirically evaluate and compare our two solutions

Example: e-commerce platform

Hierarchical database for an e-commerce platform




HTML pre-rendering



Workload-Aware Property Index

Hierarchical Database with WAPI



We denote node n 's property k as $n[k]$ and node n 's descendants as $desc(n)$.

Definition (CAS Query)

Given node m , property k and value v , a CAS query $Q(k, v, m)$ returns all descendants of m which have k set to v , i.e.,

$$Q(k, v, m) = \{n | n \in desc(m) \wedge n[k] = v\}$$

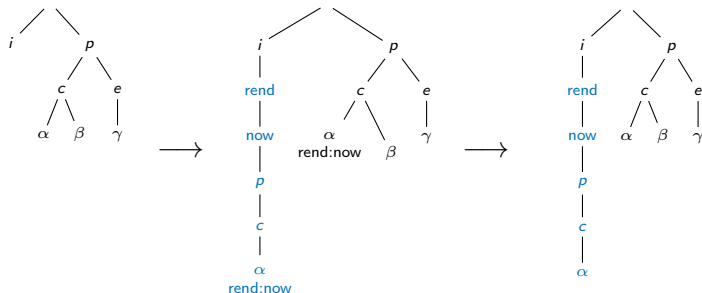
Definition (Volatile Node)

Index node n is volatile iff n 's volatility count is greater or equal than the volatility threshold τ , i.e.,

$$\text{volatile}(n) \iff \text{vol}(n) \geq \tau$$

- volatility count $\text{vol}(n)$ is the number of insertions and deletions of node n inside a sliding window.


Index nodes becoming volatile



volatile

Unproductive Nodes

Index nodes becoming unproductive



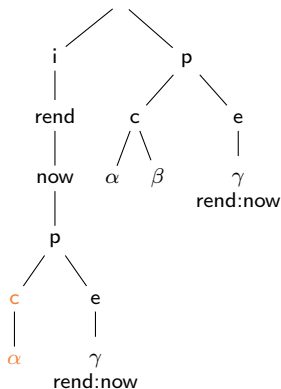
Volatile nodes might cease to be volatile and become unproductive

volatile
unproductive

Definition (Unproductive Node)

Index node n is unproductive iff n , and any descendant of n , is neither matching, nor volatile, i.e.,

$$\text{unproductive}(n) \iff \forall m (m \in (\{n\} \cup \text{desc}(n)) \implies (\neg \text{matching}(m) \wedge \neg \text{volatile}(m)))$$



$$unproductive(n) \iff \forall m(m \in (\{n\} \cup desc(n)) \implies (\neg matching(m) \wedge \neg volatile(m)))$$

volatile
 unproductive

The number of unproductive nodes depends on:

- Volatility threshold τ
- Sliding window length L
- Workload skew s
- Update operations per second

Unproductive index node cleaning, we propose:

- Periodic Garbage Collection (GC)
- Query-Time Pruning (QTP)

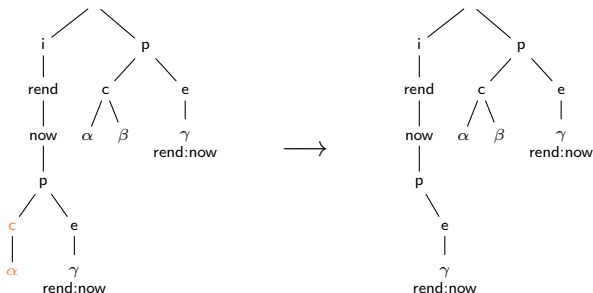
Periodic Garbage Collection (GC)

Periodic GC

Main idea:

- Background process
- Periodically traverse the whole index subtree
- Prune any visited unproductive node

Periodic GC



volatile

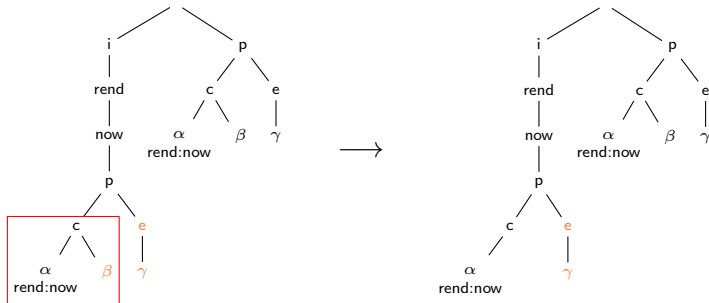
Query-Time Pruning (QTP)

Query-Time Pruning

Main idea:

- Prune unproductive nodes during query execution
- Piggybacking on query execution
- Adds overhead on query runtime
- Avoids unnecessary full index traversals
- We traverse a part of the index subtree and prune any visited unproductive node

Query-Time Pruning



unproductive

$Q(\text{render}, \text{now}, \text{/products/consumer})$

Experimental Evaluation


Experiments are run on the hierarchical database system
Apache Jackrabbit Oak




Datasets

- Datasets resemble the content subtree

(a) Synthetic



(b) Dell

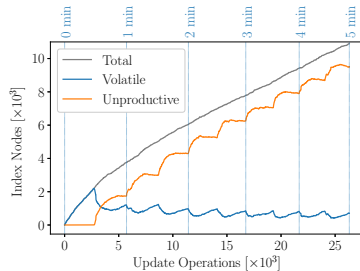
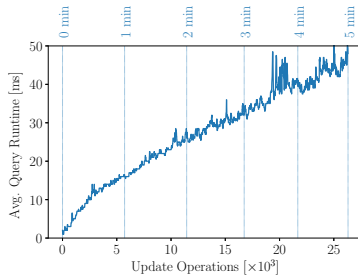


Workload simulation

- Zipf distribution
- Workload changes every 30 seconds
- 10 update operations per query operation
- Update operation: add or remove a property triggering index updates

Impact of Unproductive Nodes on Query Runtime

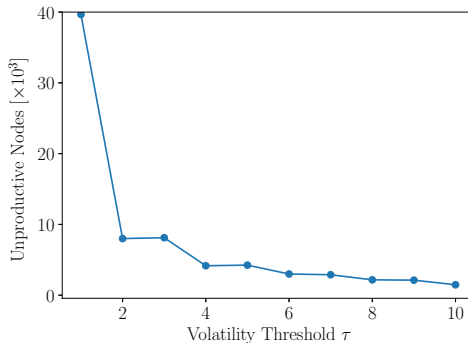
Impact of Unproductive Nodes on Query Runtime



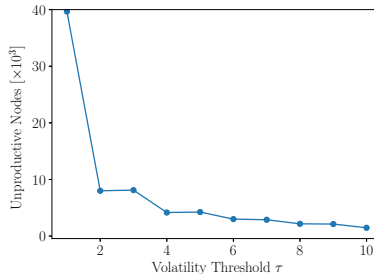
- Query runtime increases by an order of magnitude after five minutes

Volatility threshold τ

Volatility threshold τ



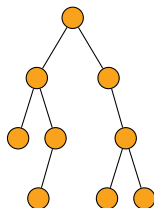
Volatility threshold τ



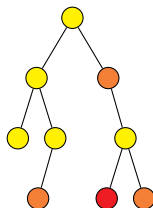
- $\tau \nearrow \implies$ volatile nodes \searrow
- volatile nodes $\searrow \implies$ unproductive nodes \searrow
- Power law relationship between #unproductive nodes and τ

Workload skew s

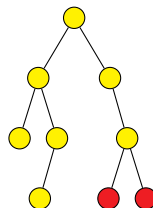
Workload skew s



No skew
 $s = 0$

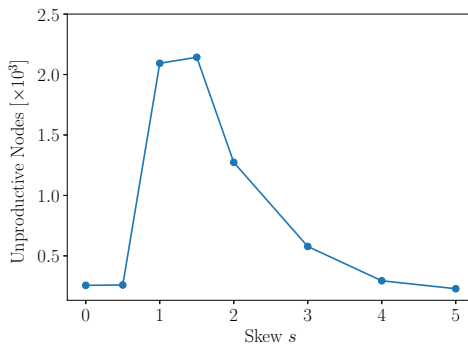


Normal skew
 $s = 1$

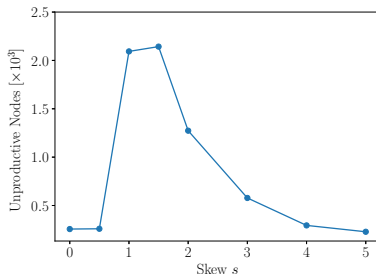


High skew
 $s = 2$

Workload skew s





Workload skew s



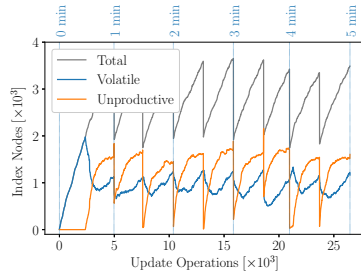
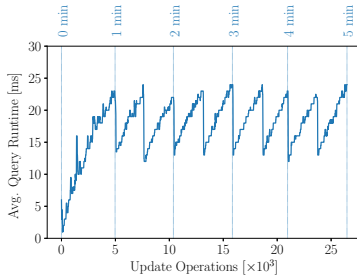
- $s > 1 \implies$ small hotspot \implies few unproductive nodes
- $s < 1$ (uniform) \implies no hotspot \implies few unproductive nodes

Cleaning Unproductive Nodes

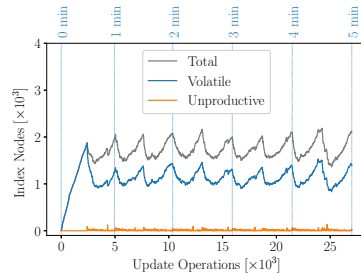
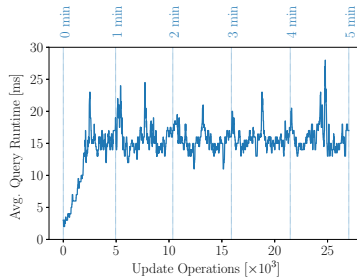


We run GC every 30 seconds

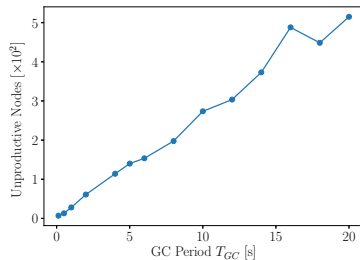
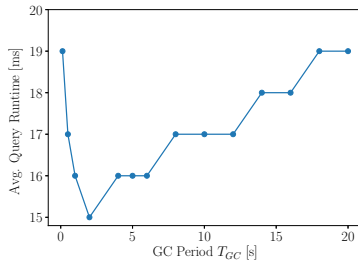
GC



QTP



GC period T_{GC}



- Optimal GC period T_{GC}^* : period with the smallest query runtime
- Too small $T_{GC} \implies$ GC steals resources from query executor

QTP

Algorithm: QueryQTP

Data: Query $Q(k, v, m)$, where k is a property, v a value and $m (= / \lambda_1 / \dots / \lambda_d)$ a content node's path.

Result: A set of nodes satisfying $Q(k, v, m)$

$r \leftarrow \emptyset$

for $node\ n \in desc(/i/k/v/\lambda_1/\dots/\lambda_d)$ in

postorder tree walk do

if $matching(n)$ then

$r \leftarrow r \cup \{ *n \}$


else if $children(n) = \emptyset \wedge \neg volatile(n)$ then

 delete node n

return r


runtime [ms]

traversal 50% -  10 -

matching 2.5% -  0.5 -

children 30% -  6 -

delete 15% -  3 -

volatile 2.5% -  0.5 -


Periodic GC vs. QTP

Simple model:

- assume constant rate of growth of unproductive nodes


Simple Model

- production rate of unproductive nodes
 $r = \frac{u}{s}$ (u unproductive nodes per second)



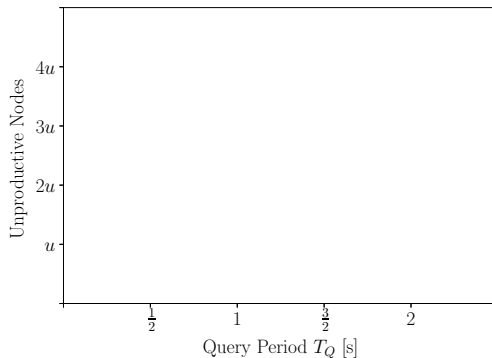
Simple Model, GC

- production rate $r = \frac{u}{s}$
- GC period $T_{GC} = 2s$



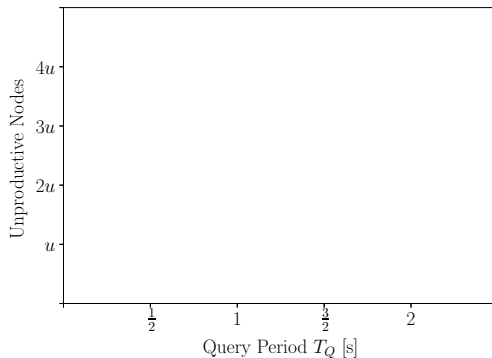
Simple Model, GC vs. QTP

- production rate $r = \frac{u}{s}$
- GC period $T_{GC} = 2s$

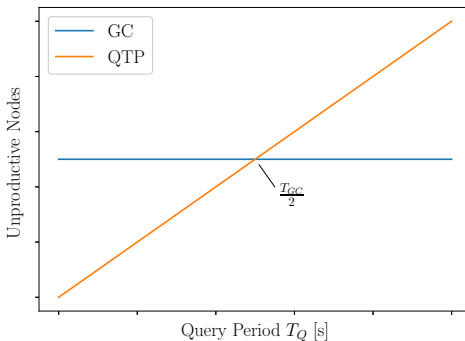


Simple Model, GC vs. QTP

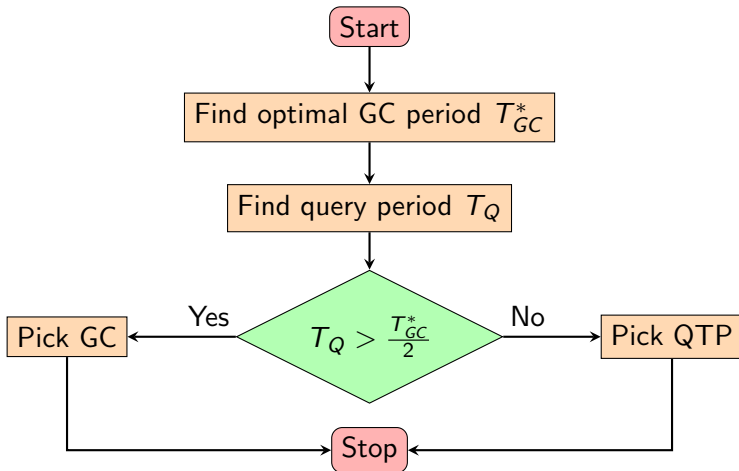
- production rate $r = \frac{u}{s}$, $r' = \frac{2u}{s}$
- GC period $T_{GC} = 2s$



Simple Model, GC vs. QTP



Queries traverse on average fewer unproductive nodes under QTP when the query period T_Q is less than half the GC period T_{GC}



Conclusion

Unproductive Nodes

- volatility threshold $\tau \nearrow \implies$ unproductive nodes \searrow
- sliding window length $L \nearrow \implies$ unproductive nodes \nearrow
- workload skew $s \updownarrow \implies$ unproductive nodes \searrow
- update operations per second $\nearrow \implies$ unproductive nodes \nearrow

GC & QTP

GC

- Periodically cleans unproductive index nodes
- Sawtooth pattern
- Slows down system if run too often

QTP

- Faster and more stable than GC when queries are frequent
- Adds overhead to queries
- Overhead negligible in the long-term

Future Work

Future Work

- Concurrency control
- Frequently changing query filter
- Unproductive node production rate