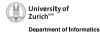
An Adaptive Index for Hierarchical Database Systems

Rafael Kallis

BSc Thesis

February 4, 2018





The Workload-Aware Property Index (WAPI):

- Detects frequently updated nodes
- Stops pruning such volatile nodes
- Significantly improves update throughput

Unproductive Nodes are an unwanted byproduct:

- When the workload changes, volatile nodes cease to be volatile
- They waste space and slow down queries
- They do not contribute to a query match and contain no data

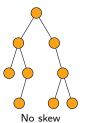
In this thesis we:

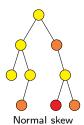
- Design and implement two solutions in order to mitigate unproductive nodes
- Analyze the factors impacting the production of unproductive nodes
- Empirically evaluate and compare our two solutions

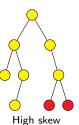
A Content Management System's (CMS) workload is:

- skewed
- update-heavy
- changing

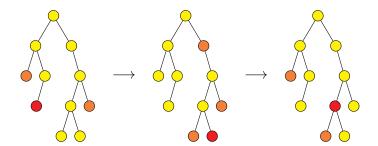
Skewed workload: small subset of nodes gets frequently updated







Changing workload: as time passes, hotspots change



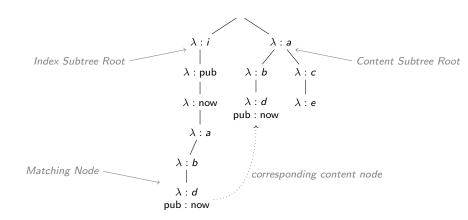
Abstract & Outline CMS Workload Workload-Aware Property Inde:

CMSs usually use a job-queuing system that has the noted characteristics

Abstract & Outline CMS Workload Workload-Aware Property Index

Workload-Aware Property Index

Hierarchical Database with WAPI



We mostly executes content-and-structure (CAS) queries [1]. We denote node n's property k as n[k] and node n's descendants as desc(n).

Definition (CAS Query)

Given node m, property k and value v, a CAS query Q(k, v, m) returns all descendants of m which have k set to v, i.e.,

$$Q(k, v, m) = \{n | n \in desc(m) \land n[k] = v\}$$

Volatility is the measure which is used by the WAPI in order to distinguish whether to remove a node or not from the index. Wellenzohn et al. [2] propose to look at the recent transactional workload to check whether a node n is volatile.

Definition (Volatility Count)

The volatility count vol(n) of index node n on database instance O_i , is the number of times node n was added or removed from snapshots contained in a Sliding Window of Length L over history H_i , i.e.,

$$vol(n) = |\{G^b | G^b \in H_i \land t(G^b) \in [t_n - L + 1, t_n] \land \exists G^a[$$

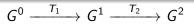
$$G^a = pre(G^b) \land ([n^a \notin N(G^a) \land n^b \in N(G^b)] \lor [n^a \in N(G^a) \land n^b \notin N(G^b)])\}|$$

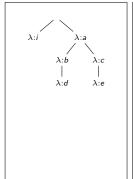
Definition (Volatile Node)

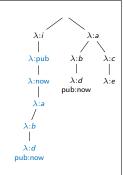
Index node n is volatile iff n's volatility count is greater or equal than the volatility threshold τ , i.e.,

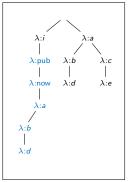
$$volatile(n) \iff vol(n) \ge \tau$$

Index nodes becoming volatile









Snapshot
$$G^0$$

 $t(G^0) = t$

Snapshot
$$G^1$$

 $t(G^1) = t + 1$

Snapshot
$$G^2$$

 $t(G^2) = t + 2$

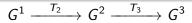
$$\tau = 1, L = 2$$

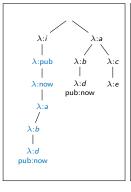
Introduction
Periodic Garbage Collection
Query-Time Pruning

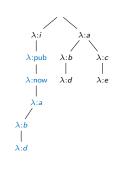
Unproductive Nodes

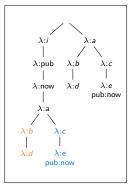
- \bullet When nodes are volatile, their volatility count has to be at least τ
- Time passes and the database workload changes
- Insertions and deletions that increased the volatility count drop out of the sliding window
- ullet If the volatility count drops below threshold au, the node ceases to be volatile
- If the same holds for the node's descendants, we call the node and its descendants unproductive

Index nodes becoming unproductive









Snapshot
$$G^1$$

 $t(G^1) = t + 1$

Snapshot
$$G^2$$

 $t(G^2) = t + 2$

Snapshot
$$G^3$$

 $t(G^3) = t + 3$

Variables impacting unproductive node production rate:

- ullet Volatility threshold au
- Sliding window length L
- Workload skew s
- Update tx per second

Unproductive index node cleaning, we propose:

- Periodic Garbage Collection (GC)
- Query-Time Pruning (QTP)

Introduction
Periodic Garbage Collection
Query-Time Pruning

Periodic Garbage Collection (GC)

Periodic GC

Main idea:

- Background process
- Periodically traverse index subtree
- Prune any visited unproductive node

Periodic GC

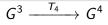
```
Algorithm: GarbageCollect
```

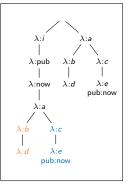
```
for node n \in desc(/i) in postorder tree walk do

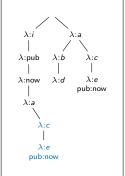
if chd(n) = \emptyset \land \neg matching(n) \land \neg volatile(n) then

delete node n
```

Periodic GC







Snapshot
$$G^3$$

 $t(G^3) = t + 3$

Snapshot
$$G^4$$

 $t(G^4) = t + 4$

Introduction
Periodic Garbage Collection
Query-Time Pruning

Query-Time Pruning (QTP)

Query-Time Pruning

Main idea:

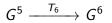
- Prune unproductive nodes during query execution
- Piggypacking on query execution
- Adds overhead on query runtime

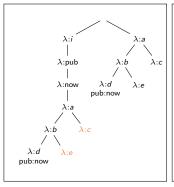
Query-Time Pruning

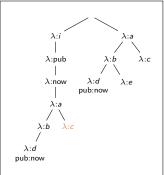
return r

```
Algorithm: QueryQTP
Data: Query Q(k, v, m), where k is a property, v a value and
        m (= /\lambda_1/.../\lambda_d) a content node's path.
Result: A set of nodes satisfying Q(k, v, m)
r \longleftarrow \emptyset
for node n \in desc(/i/k/v/\lambda_1/.../\lambda_d) in postorder tree walk
 do
    if matching(n) then
    r \longleftarrow r \cup \{*n\}
    else if chd(n) = \emptyset \land \neg volatile(n) then
     ∟ delete node n
```

Query-Time Pruning







Snapshot G⁵

Snapshot G⁶

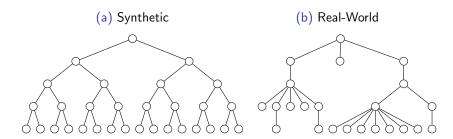
 $\tau = 1, L = 2, Q(pub, now, /a/b)$



Unproductive Nodes Periodic Garbage Collection Query-Time Pruning Comparison

Experimental Evaluation

Datasets



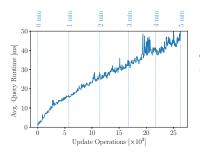
Workload simulation

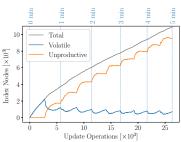
- Zipf distribution
- Changes every 30 seconds
- 10 update operations per query operation
- Update operation: add and remove an index node

Unproductive Nodes Periodic Garbage Collection Query-Time Pruning Comparison

Impact of Unproductive Nodes on Query Runtime

Impact of Unproductive Nodes on Query Runtime

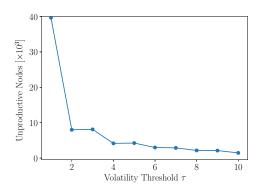




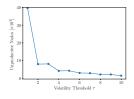
Introduction Unproductive Nodes Experimental Evaluation Conclusion Unproductive Nodes Periodic Garbage Collection Query-Time Pruning Comparison

Volatility threshold au

Volatility threshold au



Volatility threshold au

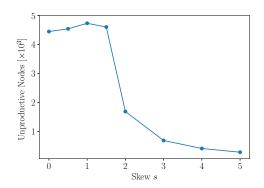


- ullet au increases \Longrightarrow nodes are less likely to become volatile
- Fewer volatile nodes ⇒ fewer unproductive nodes
- ullet Power law relationship between #unproductive nodes and au

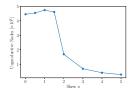
Introduction Unproductive Nodes Experimental Evaluation Conclusion Unproductive Nodes
Periodic Garbage Collection
Query-Time Pruning
Comparison

Workload skew s

Workload skew s



Workload skew s

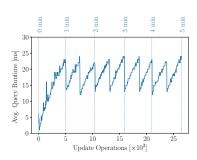


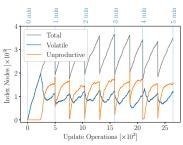
- very high $s \implies$ very small hotspot
- very small s (uniform) \Longrightarrow no hotspot

Unproductive Nodes Periodic Garbage Collection Query-Time Pruning Comparison

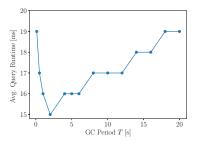
Periodic Garbage Collection

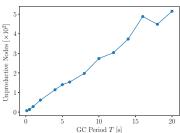
Periodic GC





GC period T

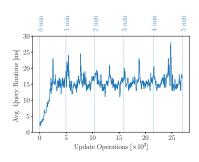


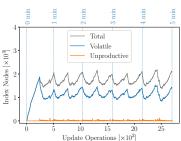


Unproductive Nodes Periodic Garbage Collection Query-Time Pruning Comparison

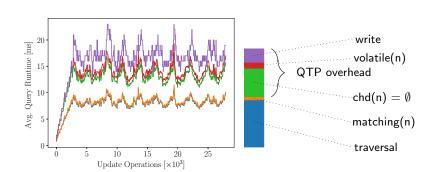
Query-Time Pruning

QTP





QTP



QTP

```
Algorithm: QueryQTP
```

```
Data: Query Q(k, v, m), where k is a property, v a value and m (= /\lambda_1 / ... / \lambda_d) a content node's path.
```

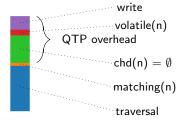
```
Result: A set of nodes satisfying Q(k, v, m) r \leftarrow \emptyset for node n \in desc(/i/k/v/\lambda_1/.../\lambda_d) in postorder tree walk do

if matching(n) then

r \leftarrow r \cup \{*n\}
else if chd(n) = \emptyset \land \neg volatile(n) then
```

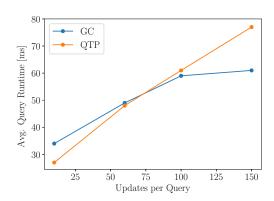
delete node n

return r



Unproductive Nodes Periodic Garbage Collection Query-Time Pruning Comparison

Periodic GC Vs. QTP



References



MATHIS, C., HÄRDER, T., SCHMIDT, K., AND BÄCHLE, S. XML indexing and storage: fulfilling the wish list. Computer Science - R&D 30, 1 (2015), 51–68.



Wellenzohn, K., Böhlen, M., Helmer, S., Reutegger, M., and Sakr, S. A Workload-Aware Index for Tree-Structured Data. To be published.